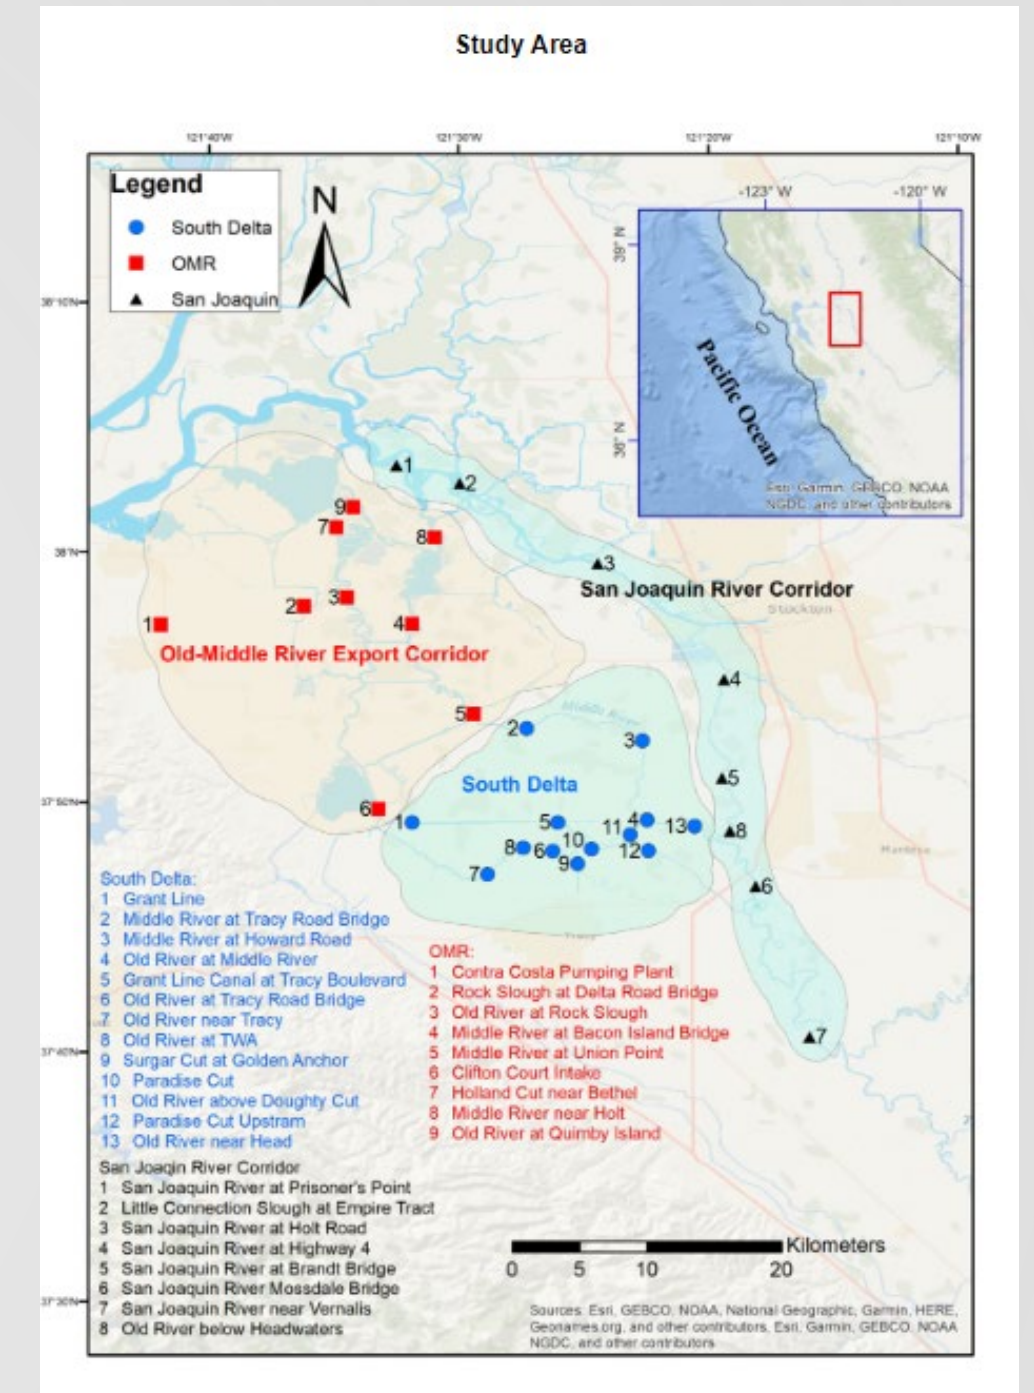


Hands-on

Goal: Build and evaluate ML models to predict ion concentrations in the Interior Delta

What will be covered:

1. **Data Loading:** Import water quality datasets
2. **Data Visualization:** Explore relationships between variables
3. **Preprocessing:** Feature engineering, scaling, and encoding
4. **Model Training:** Train multiple ML models
5. **Model Evaluation:** Compare performance across models
6. **Interactive Dashboard:** Create a user-friendly interface



Disclaimer: This notebook is created for educational purposes as part of a machine learning workshop. The models and methods demonstrated here are simplified examples and not necessarily optimized for production use. The actual models used in production environments may incorporate additional features, more sophisticated techniques, and more extensive validation.



Data Preprocessing

Scaling, and encoding

Predictors:

1- EC

2- Sacramento X2

3- Month

4- Water Year Type

5- Location

Numerical values: feature scaling


Prevent features with larger magnitudes from dominating those with smaller magnitudes

Categorical values: hot encoding


Enable machine learning models to interpret and utilize non-numeric data effectively

Before

After



WYT
W
D
C
AN
BN



W	D	C	AN	BN
1	0	0	0	0
0	1	0	0	0
0	0	1	0	0
0	0	0	1	0
0	0	0	0	1

Binary encoded



CALIFORNIA DEPARTMENT OF
WATER RESOURCES


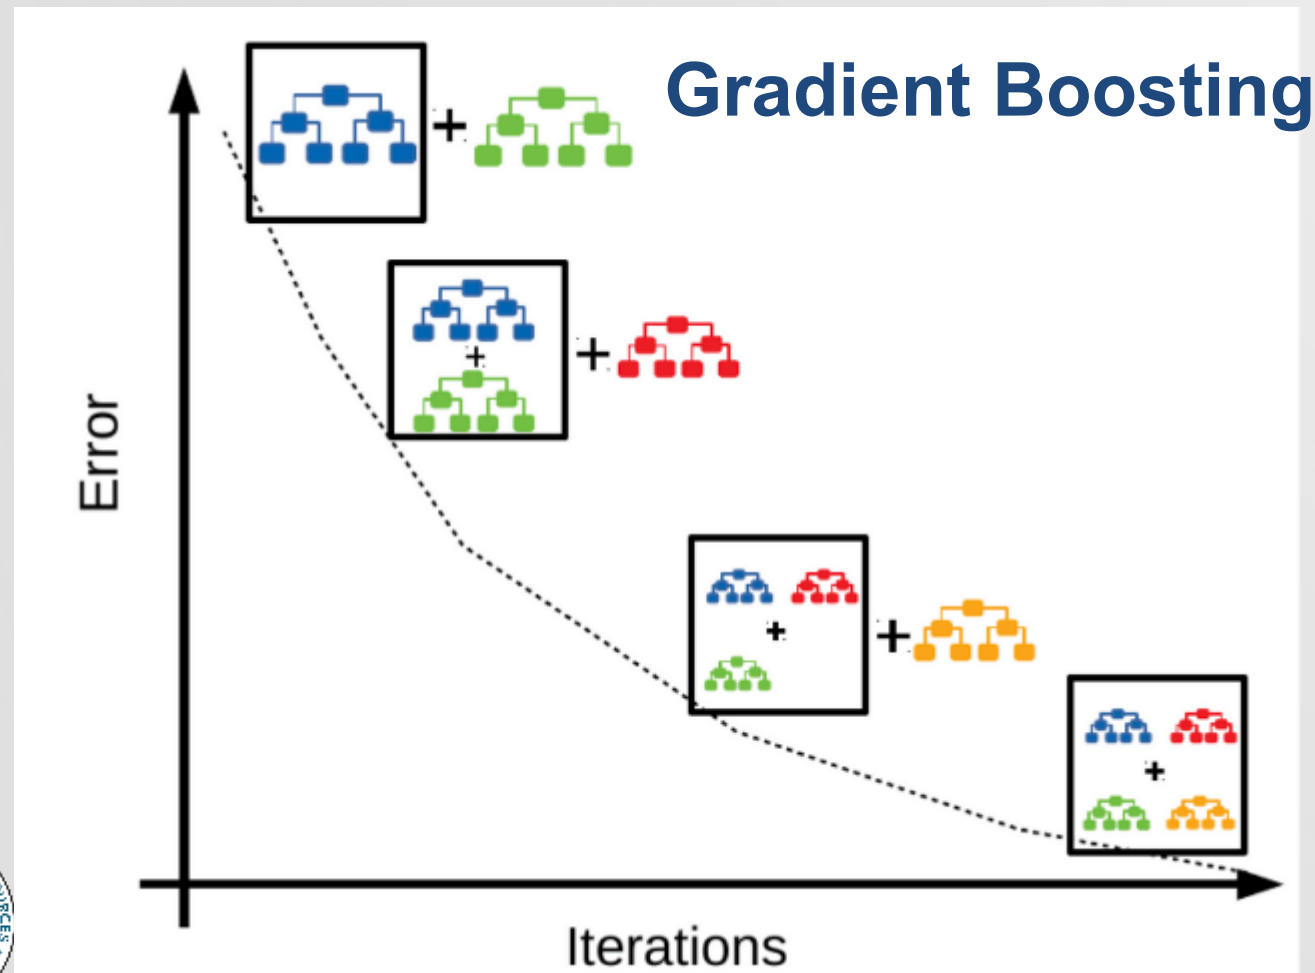
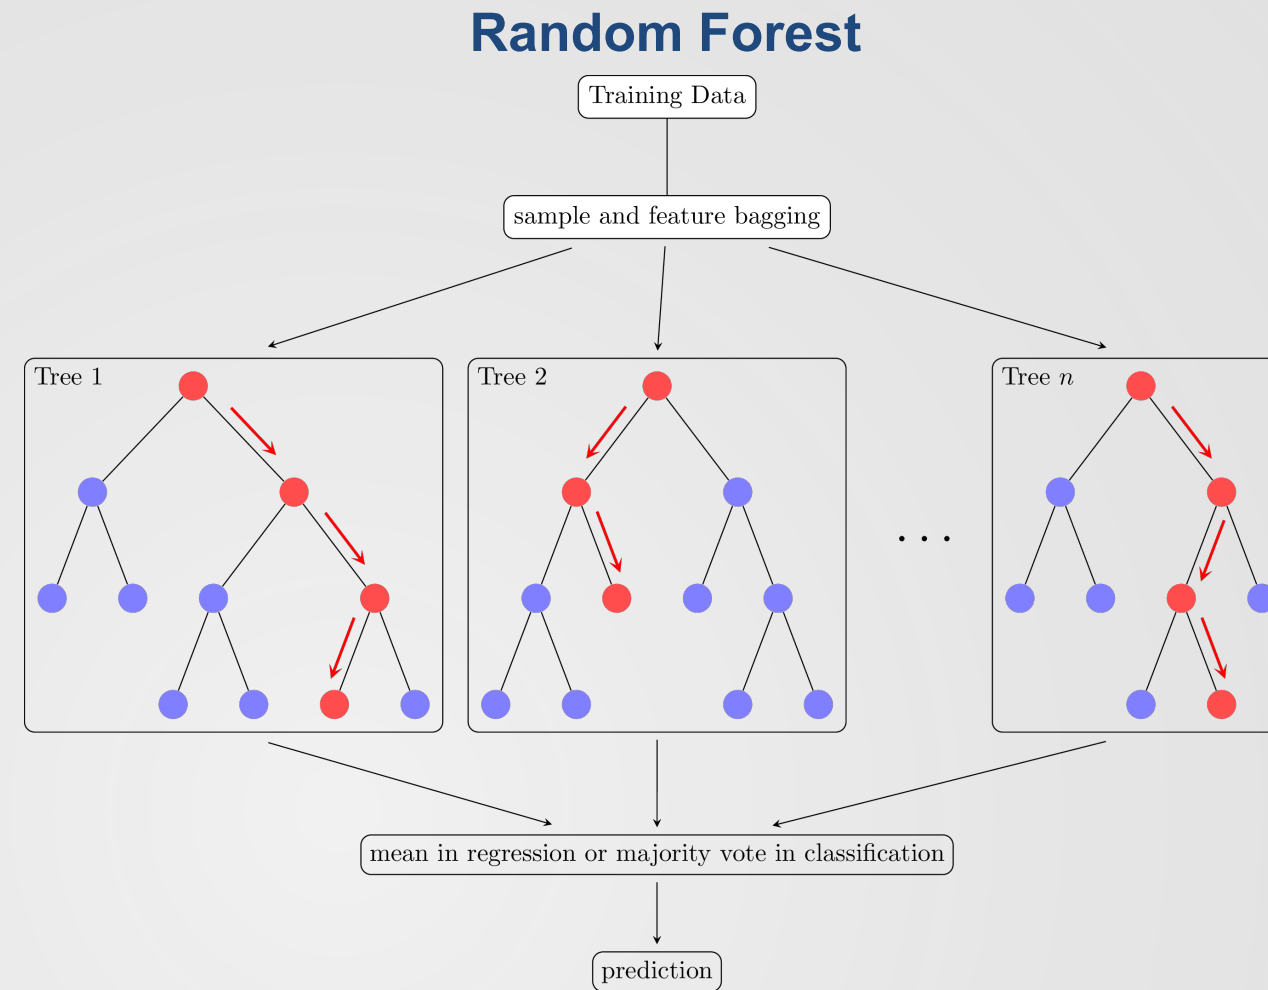
Model Selection

Models:

➤ Ensemble method:

- Random Forest (RF)
- Gradient Boosting (GB)

➤ Artificial Neural Networks (ANN)

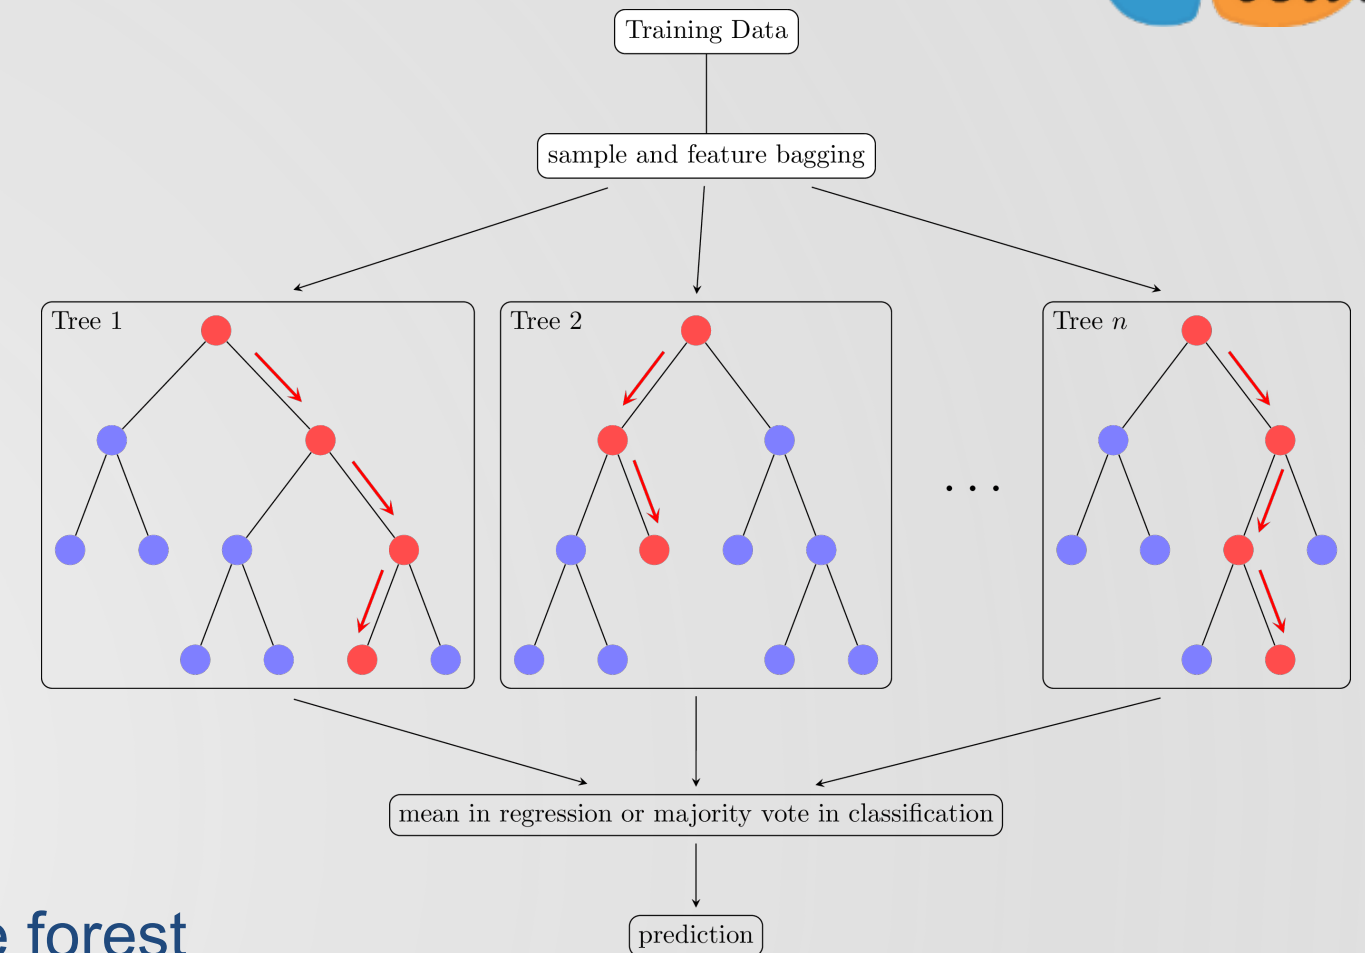


Random Forest (RF)



Key Characteristics:

- Trees built in parallel (independent)
- Random subset of data for each tree (bootstrapping)
- Random subset of features at each split
- Equal weight for all trees



```
# Random Forest model
rf_model = RandomForestRegressor(
    n_estimators=100,
    max_depth=None,
    min_samples_split=2,
    min_samples_leaf=1,
    random_state=42,
    n_jobs=-1
)
```

The number of trees in the forest

The maximum depth of the tree.

The minimum number of samples required to split an internal node

The minimum number of samples required to be at a leaf node

Means using all processors



Gradient Boosting (GB)



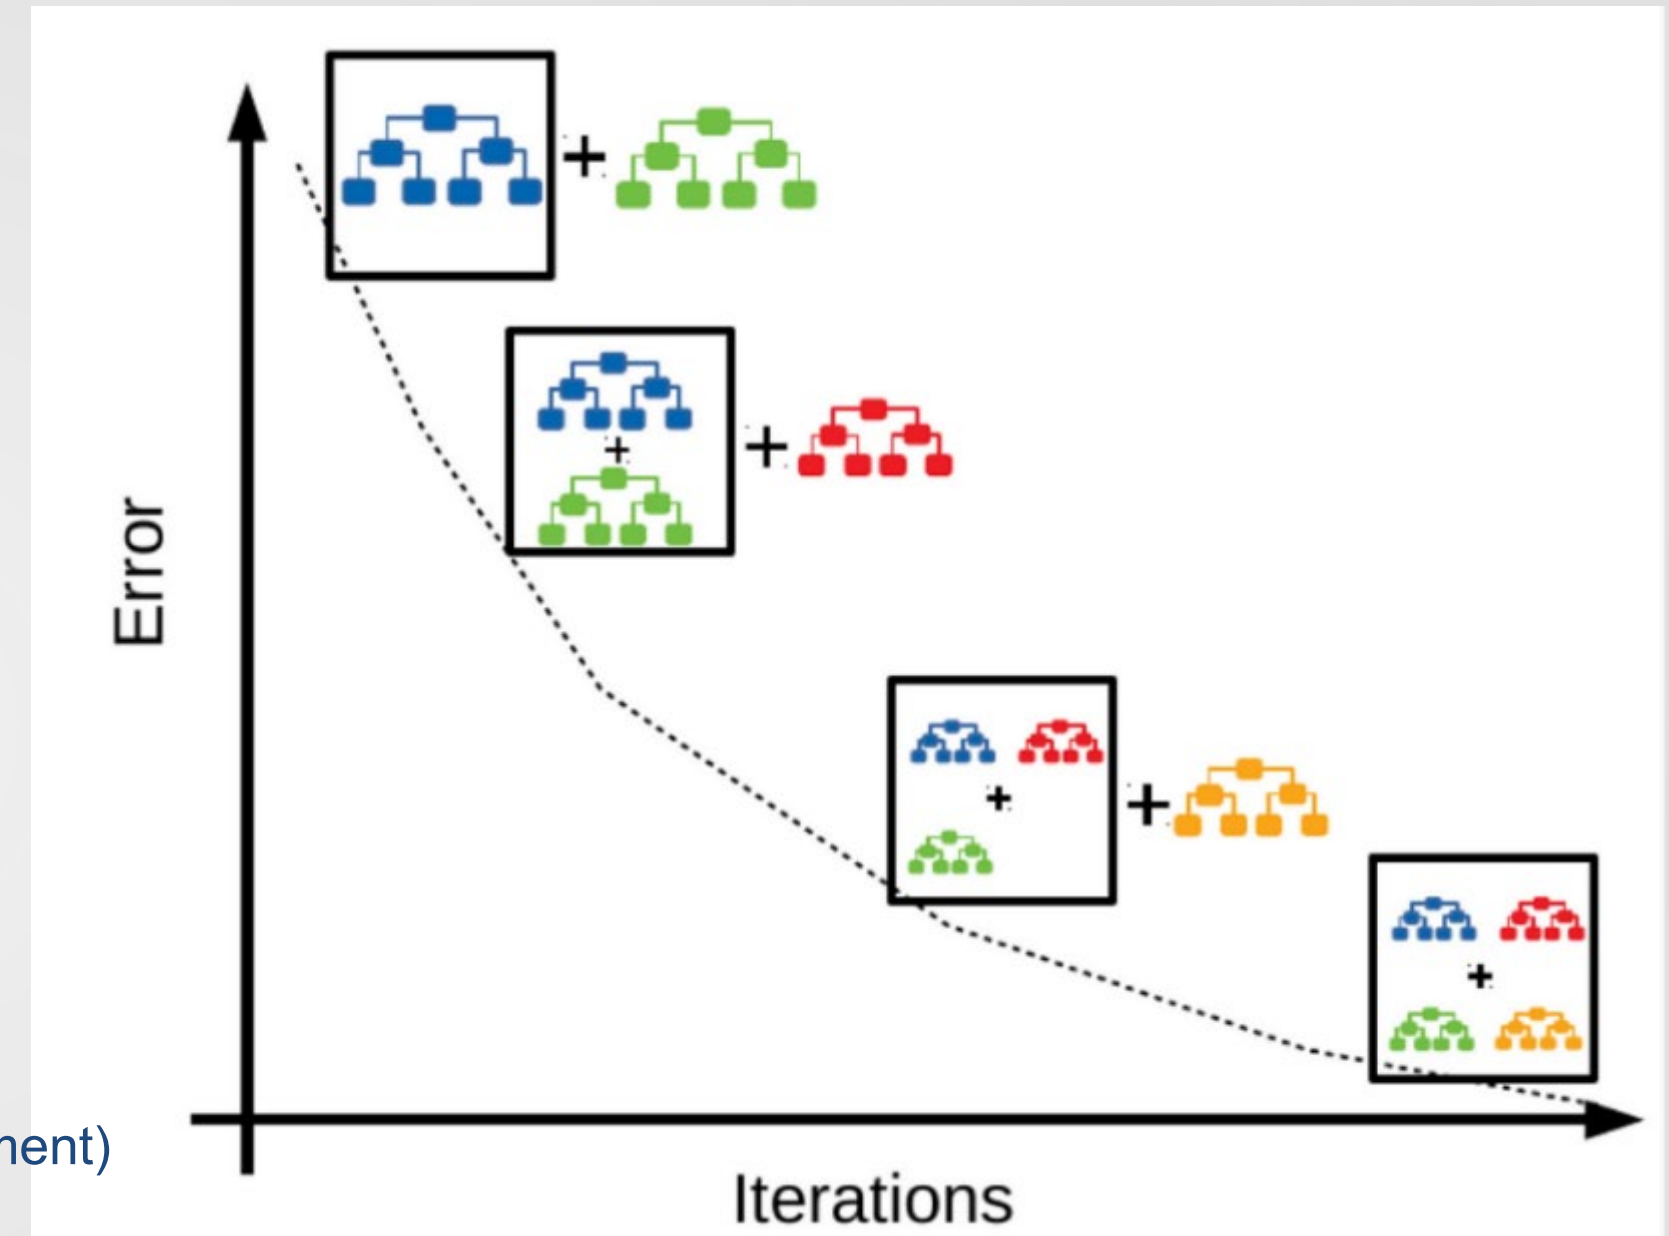
Key Characteristics:

- Trees built sequentially (dependent)
- Each tree focuses on previous errors
- Weighted combination of trees
- Learning rate controls contribution of each tree

```
# XGBoost model
xgb_model = XGBRegressor(
    n_estimators=100,
    learning_rate=0.1,
    max_depth=5,
    subsample=0.8,
    colsample_bytree=0.8,
    random_state=42,
    n_jobs=-1
)
```

Subsampling (without replacement)

Means using all processors



Artificial Neural Networks (ANN)




```
# network model
nn_model = Sequential([
    Dense(16, activation='relu', input_shape=(num_features,)),
    Dense(8, activation='relu'),
    Dense(4, activation='relu'),
    Dense(1)
])

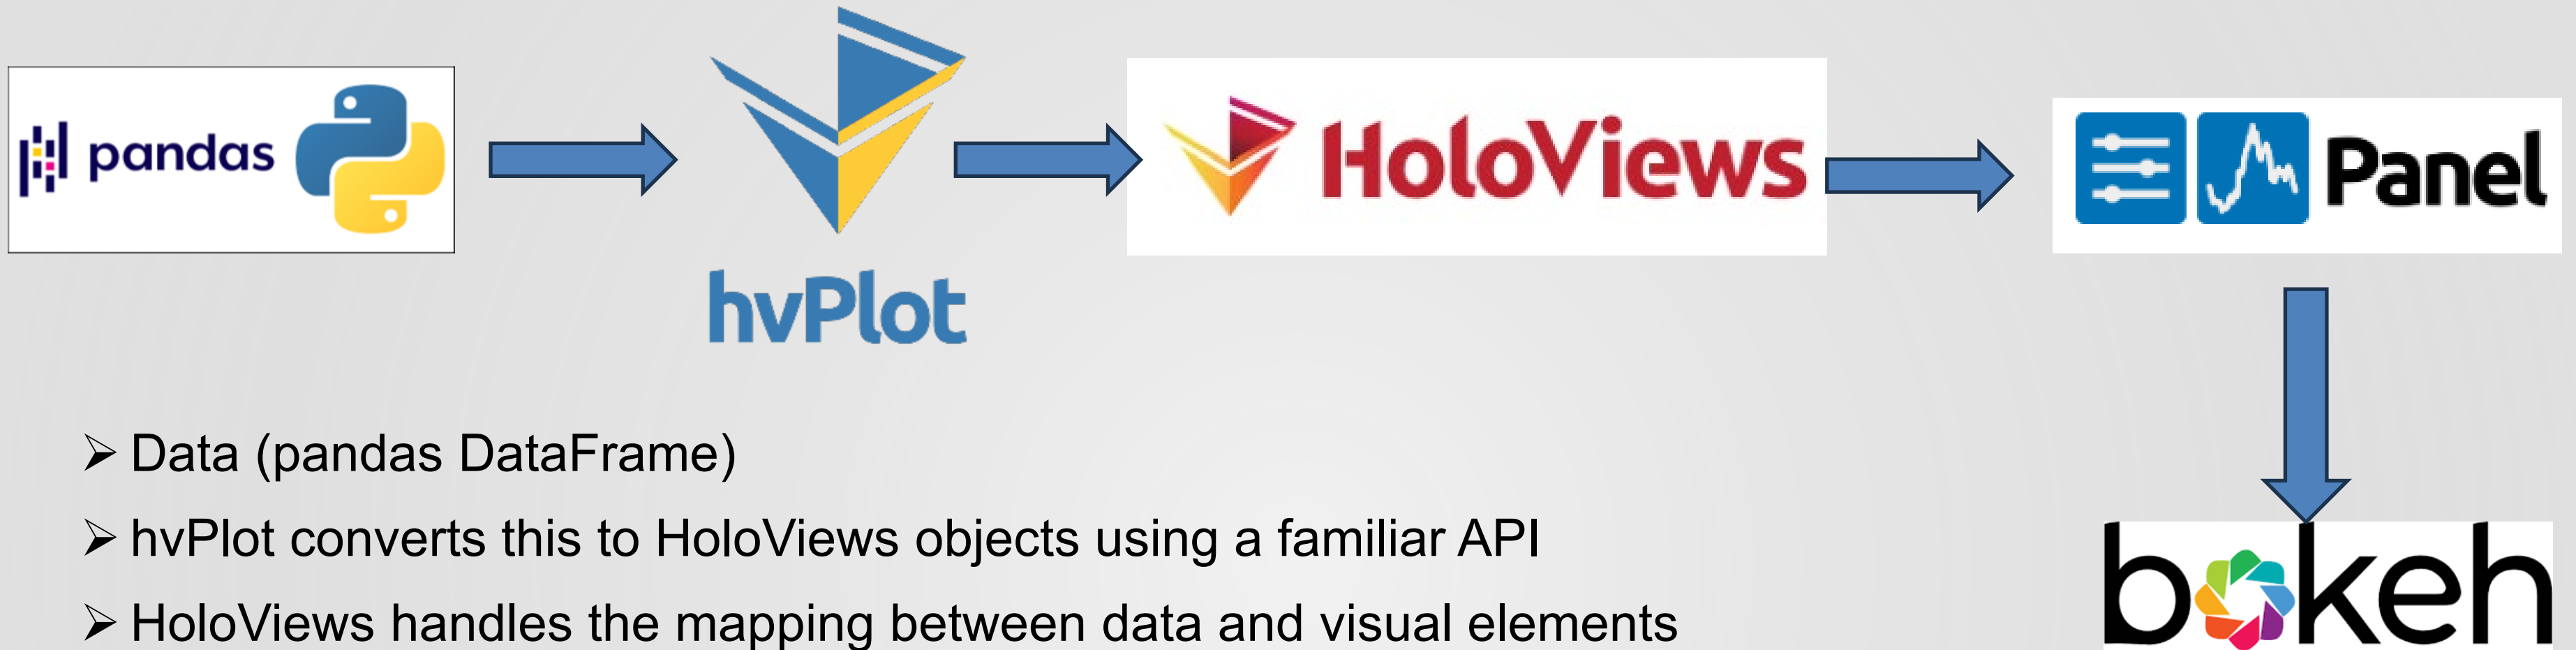
# Compile the model
nn_model.compile(
    optimizer=tf.keras.optimizers.Adam(learning_rate=0.001),
    loss='mean_squared_error'
)

# Early stopping to prevent overfitting
early_stopping = EarlyStopping(
    monitor='val_loss',
    patience=20,
    restore_best_weights=True,
    verbose=1
)

# Train the model
nn_model.fit(
    X_train_processed_filtered,
    y_train.values.astype(np.float32),
    epochs=100,
    batch_size=32,
    validation_split=0.2,
    callbacks=[early_stopping],
    verbose=1
)
```



Dashboard



- Data (pandas DataFrame)
- hvPlot converts this to HoloViews objects using a familiar API
- HoloViews handles the mapping between data and visual elements
- Bokeh (as the backend renderer) creates the actual JavaScript-based visualization

(<https://docs.bokeh.org/en/latest/docs/gallery.html>)

- Panel arranges these visualizations and adds interactive controls

(<https://panel.holoviz.org/reference/index.html#widgets>)



CALIFORNIA DEPARTMENT OF
WATER RESOURCES

Bokeh: Low-level JavaScript-based plotting (handles the actual rendering)
HoloViews: Mid-level data-centric visualization objects
hvPlot: High-level pandas-like plotting API
Panel: Dashboard composition and widget integration


Practice Problem 1: Simulate SO4 Concentration

Objective: Repeat the modeling process we completed for Bromine (Br) but now for Sulfate (SO4).

Step 1: In the Part 3 (Preprocessing) Change `TARGET_ION = 'Br'` to `TARGET_ION = 'SO4'`

Step 2: Run all code cells from Section 3 (Preprocessing) to the end

Step 3: Review the results and identify which model performs best for SO4 prediction



Practice Problem 2: Hyperparameter adjustment

➤ Random Forest model:


- Limit tree depth to 4 (previously unlimited)

➤ XGBoost model:

- Reduce maximum tree depth from 5 to 4

➤ Neural Network model:

- Make the network bigger: (32→16→8 neurons instead of 16→8→4)
- Train for more epochs (200 instead of 100)
- Increase patience (30 instead of 20)



Practice Problem 3: Custom Loss Function for Neural Network

Objective: Implement and evaluate a simpler custom loss function for the Neural Network model. RMSE instead of MSE. We want to do it with custom loss function.

4. Model Training

```
# ===== MODEL 3: NEURAL NETWORK =====
print(f"\nTraining Neural Network for {TARGET_ION}...")

# Get the number of input features from the filtered data
num_features = X_train_processed_filtered.shape[1]

# network model
nn_model = Sequential([
    Dense(32, activation='relu', input_shape=(num_features,)),
    Dense(16, activation='relu'),
    Dense(8, activation='relu'),
    Dense(1)
])

# Compile the model
nn_model.compile(
    optimizer=tf.keras.optimizers.Adam(learning_rate=0.001),
    loss='mean_squared_error'
)
```



```
# ===== MODEL 3: NEURAL NETWORK =====
print(f"\nTraining Neural Network for {TARGET_ION}...")

def rmse_loss(y_true, y_pred):
    """
    RMSE = sqrt(mean((y_true - y_pred)²))
    """
    y_true = tf.cast(y_true, tf.float32)
    y_pred = tf.cast(y_pred, tf.float32)
    squared_error = tf.square(y_true - y_pred)
    mean_squared_error = tf.reduce_mean(squared_error)
    return tf.sqrt(mean_squared_error)

# Get the number of input features from the filtered data
num_features = X_train_processed_filtered.shape[1]

# network model
nn_model = Sequential([
    Dense(32, activation='relu', input_shape=(num_features,)),
    Dense(16, activation='relu'),
    Dense(8, activation='relu'),
    Dense(1)
])

# Compile the model
nn_model.compile(
    optimizer=tf.keras.optimizers.Adam(learning_rate=0.001),
    loss=rmse_loss
)
```

Step 1



Step 2



CALIFORNIA DEPARTMENT OF
WATER RESOURCES

5. Evaluation

```
# ===== MODEL EVALUATION =====

# Load the preprocessor, feature names, and models
preprocessor = joblib.load(f'{output_folder}/{TARGET_ION}_preprocessor.joblib')
feature_names = joblib.load(f'{output_folder}/{TARGET_ION}_model_feature_names.joblib')

# Load the trained models
rf_model = joblib.load(f'{output_folder}/{TARGET_ION}_random_forest_model.joblib')
xgb_model = joblib.load(f'{output_folder}/{TARGET_ION}_xgboost_model.joblib')
nn_model = load_model(f'{output_folder}/{TARGET_ION}_nn_model.h5', custom_objects={'rmse_loss': rmse_loss})
```

← Step 3

	Model	Dataset	MAE	Percentage Bias	R2
4	Neural Network	Training	9.06	9.01	0.92
5	Neural Network	Test	10.07	6.50	0.91

← Loss=MSE

4	Neural Network	Training	8.31	4.92	0.92
5	Neural Network	Test	9.73	2.49	0.91

← Loss=RMSE




```
def rmse_loss(y_true, y_pred):  
    """  
    RMSE = sqrt(mean((y_true - y_pred)2))  
    """  
    y_true = tf.cast(y_true, tf.float32)  
    y_pred = tf.cast(y_pred, tf.float32)  
    squared_error = tf.square(y_true - y_pred)  
    mean_squared_error = tf.reduce_mean(squared_error)  
    return tf.sqrt(mean_squared_error)
```



Questions?

